

# Seeing the Light with Refactoring and Testing

Richie Morrisroe

June 14, 2020

# Data Scientists and Code

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- An awful lot of us produce our results/analyses through code
- Much of this is ad-hoc investigations
- Often these investigations produce output which others want
- You create a cron job
- Congratulations, you are now supporting production code <sup>1</sup>

---

<sup>1</sup>in the sense that others rely on it. SQL tables are often the worst offenders (as they are often the easiest to create)

# My background

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- I trained as a psychologist
- I like numbers
- So I ended up as a data scientist
- At no point have I ever had a class on coding
- I haven't even had that many classes on statistics

# So why are you even doing this talk?

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- Because I have learned, from bitter, bitter experience that this stuff is important
- People (hopefully) make decisions as a result of our analyses
- Often, these decisions can have far-ranging impacts, most of which are impossible <sup>2</sup> to foresee
- As professionals, we have an obligation to make sure these results are correct
- This means we need to make sure our code is correct

---

<sup>2</sup>i.e. hard

# This Talk

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- I will talk about tools to improve the quality of code, and our confidence in it
- These tools are:
  - Automated testing
  - Refactoring
  - Test-Driven Development

# Before Refactoring

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- You **must** write tests
- Otherwise you <sup>3</sup> **will** introduce bugs

---

<sup>3</sup>or at least me

# Types of Testing

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- Individual functions: unit testing
- Processes working together: integration testing
- Overall application: functional testing
- Lots of space between these points

# My Code is so awful that it can't be tested

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- This is depressingly normal
- You want some kind of characteristic test
- This is often a dataframe type structure
- The quickest test is to ensure that both dataframes are equal
- There is a *wonderful* book that guides you through this



# Working Effectively with Legacy Code

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- The book, by Michael Feathers is really useful
- Legacy code = code without tests
- Book is based on particular problems
- And provides approaches for solving them
- Catalog of refactorings that can (theoretically) be done without tests <sup>4</sup>

---

<sup>4</sup>you **always** want at least some characterisation tests (i.e. output)

# Characterisation Tests

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
require(testthat)
output_old <- readr::read_csv(
  "older_data.csv",
  ## you'll need this argument
  col_types=NULL)
test_that('output old = output new',
  all.equal(output_old, output_new))
```

- capture how the system behaves right now, and make sure it doesn't change
- This simple test can allow you to make a *lot* of progress relatively quickly
- these tests are useful to help build structure and get other tests in place
- but they prevent you from improving the code in pretty much any way

# Seams

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- This is one of the most interesting parts of WELC, which defines a seam as a place where you can alter what code runs without changing any code
- For instance, if you have some kind of folder with different versions of code then you can create a seam by changing which folder is pointed to
- This is also a very quick way of getting characterisation tests
- The simplest example is mocking out a package/script file/module by changing the path

# Pinch Points

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- Another concept from WELC
- Refers to the point where you can test as much as possible of functionalities
- Normally a dataframe of some description (input data, output data, predictions etc)

# First Steps

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- Gather the source files you are interested in
- Run them, make sure they work.
- DO NOT SKIP THIS STEP!!!

# My Example

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- Its based on the irish property market data
- I have been working on and off on this for a number of years
- I have multiple folders, datasets and scripts that perform data loading, processing and modelling
- I now want to turn this into an interactive app, and keep it updated
- I aim to add tests and refactor to make it easier for me to build on top of this foundation

# First Practical Step

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
require(devtools)
if (!dir.exists("ppr")) {
  usethis::create_package("ppr")
}
```

```
usethis::use_package("dplyr")
usethis::use_package("rlang")
usethis::use_package("readxl")
usethis::use_package("caret")
usethis::use_package("rgdal")
usethis::use_package("sp")
usethis::use_package("glmnet")
usethis::use_package("sf")
usethis::use_package("vtreat")
usethis::use_data_raw()
usethis::use_pipe()
```

# Why a package?

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- Lots of good things:
  - automated dependency management
  - easy documenting facilities (with roxygen2)
  - easy to run tests (with testthat)
  - check is just a bunch of better programming practices
- Bad things:
  - more effort
  - appeasing the mighty gods of check takes time



# Next steps

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- We have three scripts:
  - `prep_modelling.R`
  - `feature_engineering.R`
  - `models.R`

# Step 0: Create a custom environment

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- I use conda for this
- It supports multiple languages, and compiled dependencies
- This is pretty awesome for python, and its a little more light-weight than Docker
- It supports R (and apparently loads of other languages too)

```
conda create -n PPR-r r=3.6.3
```

```
conda env export --name PPR-r > PPR-r.yml
```

# First, generate or create ground truth data

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- Our scripts need to be run in order
- So, we can use the boundaries between them as points where we can introduce test points
- However, our scripts suck and therefore we don't have any convenient boundaries available

```
rm(list=ls())  
source("scripts/prep_modelling.R")
```

- Simplest way to run script
- We can then look at what it produces and decide what to test
- Biggest benefit here is not having to change anything

# Assessing what we need

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

ls()

- We have a bunch of dataframes, some functions and a few intermediate results
- However, the best **pinch point** is ppr\_train and test, as they are downstream of all our feature engineering

# Add data to package

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
ppr <- readxl::read_excel("~/Dropbox/PPR/PPR-ALL.xlsx")
set.seed(49)
ppr_sample <- dplyr::sample_frac(ppr, size = 0.05)
ppr <- ppr_sample
usethis::use_data(ppr, overwrite = TRUE)
```

- This is important for running the tests
- You *can* use absolute paths and stuff, but it's pretty ugly
- CRAN strongly suggests that data be less than 5Mb, so we take a 5% sample

# Create pinch point

- We need to run our script with the package data, so we can keep everything consistent
- We add all the scripts to inst/, which is where we'll put the test data also

```
readr::write_csv(ppr_train2,  
                 path = "prep_modelling_output_old.csv"
```

- We put this at the end of our script
- Then, we run it like so:

```
cd ppr/inst/  
Rscript prep_modelling.R
```

- We then output the full file
- However, because we took a 5% sample earlier, we need to change things

# Handling the sample

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe


- We have the data in our package <sup>5</sup>
- We need to change our script to use the package data
- We can either rename the script, or rely on version control

```
cp prep_modelling.R prep_modelling_refactor.R
```

- Then we rename our output file

```
## readr::write_csv(ppr_train2,  
## path = "prep_modelling_output_old.csv")  
readr::write_csv(ppr_train2,  
                  path = "prep_modelling_output_refacto
```

---

<sup>5</sup>right now, it's the *only* thing in the package 

# Commented out code is bad, right?

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- In general commented out code is a massive anti-pattern
- But right now, we are making minimal changes to ensure that we can safely update our script



# Import data from package

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
## library(tidyverse)
## library(readxl)
## ppr <- read_excel("~/Dropbox/PPR/PPR-ALL.xlsx",
##                    sheet = "PPR-ALL")
require(ppr)
data(ppr)
names(ppr)
```

# Add Test Directories

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
mkdir ppr/tests  
mkdir ppr/tests/testthat/  
touch ppr/tests/testthat/test_first.R  
touch ppr/tests/testthat/test_integration_prep_modelli
```

- You can do this with `usethis`, but it was borked for me when I tried, so I did it manually

# Integration Test

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
context("integration test prep_modelling")
old <- readr::read_csv("../../inst/prep_modelling_output.csv")
new <- readr::read_csv("../../inst/prep_modelling_output.csv")
test_that(
  "data.frame outputs are equal",
  expect_equal(old, new)
)
```

# Run the script again with the changes we've made

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- Commit your code before and after!
- We actually needed to make the data changes *first* as otherwise our "old" data wouldn't match the new(er) data

# Create a small script to run the integration tests

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- This talk is all about **automated** testing, so we'll write a script

```
system("Rscript prep_modelling.R")
testthat::test_file(
  "../tests/testthat/test_integration_prep_m
```

# Move existing functions to file

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- Move the functions we have to a file
- Change our script to load the functions from this file
- We now have a package with code in it
- Now we run our integration test again
- Next step is to add some unit tests

# Write some tests

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
require(testthat)
context("load data")
dat <- load_data("~/Dropbox/PPR/PPR-ALL.xlsx")
test_that('load_data returns a tibble',
  {expect_equal(
    class(dat)[1], "tbl_df")})
```

- The first thing we do is write a `load_data` function
- this allows us to later abstract how/where we load the data from without needing to change (as much code)
- Note that this function doesn't exist yet
- Let's write it

# Loading Data

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
load_data <- function(path) {  
  ppr <- readxl::read_excel(path, sheet = "PPR-ALL")  
  return(ppr)  
}
```

- This wouldn't normally make a whole lot of sense
- But I know that I am probably going to move to SQL storage with this project
- So centralising the data loading does make sense <sup>6</sup>

---

<sup>6</sup>given my use-cases



# Test Names

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
finished_names <-  
  c("date_of_sale_dd_mm_yyyy",  
    "address", "postal_code", "county",  
    "price", "not_full_market_price", "vat_exclusive",  
    "description_of_property",  
    "property_size_description") %>%  
  as.character()  
  
test_that('normalise_names works',  
{expect_equal(finished_names,  
              names(normalise_names(dat)))  
})
```

- We just grab the output of `normalise_names`, and test that
- this is a crummy test; try to test the actual invariants on edge cases

# Test Fix Price

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
test_that('fix price returns the correct price',  
          expect_equal(3e6, fix_price("3,000,000")))
```

- All good
- However, I think I spotted a bug in `fix_price`

# Fix Price Bug

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
fix_price <- function(x) {  
  nopunct <- gsub(",|\\\\.", "", x = x)  
  nums <- as.numeric(  
    iconv(nopunct, "latin1",  
          "ASCII",  
          sub = ""  
        )  
  )  
}
```

- Can you spot it?

# Test Case for Bug

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
test_that('fix price returns the correct price',  
          expect_equal(3000000.01,  
                        fix_price("3,000,000.01")))
```

- We actually shouldn't strip the dots, as they are valid in numbers
- We have a failing test, so we can fix this code (and make sure it stays fixed, which is normally more valuable)

# Fix Price Fix

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
fix_price <- function(x) {  
  nopunct <- gsub(",", "", x = x)  
  nums <- as.numeric(  
    iconv(nopunct, "latin1",  
          "ASCII",  
          sub = ""  
        )  
  )  
}
```

- We just remove the part of the gsub call that matches dots
- Now all our tests pass

# Generate Bootstrap Results

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- I'll leave testing this for now, as I probably need to do the feature normalisation first

# Break out components

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- You want to break out chunks of code into manageable (well-named) sizes
- This is the extract function refactoring (probably the most common)

```
ppr3 <- mutate(ppr2,  
               price = price / 100,  
               is_big = ifelse(price >= 2e6,  
                               "Big",  
                               "Not Big"))
```

# Write test for new function

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
test_df_output  <- readr::read_csv("../../inst/mark_val
test_df_input   <- readr::read_csv("../../inst/mark_val
test_that('mark values as large works', {
  expect_equal(mark_values_as_large(test_df_input,
                                    large = 1e6),
               test_df_output)
})
```



# Write the actual function

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
mark_values_as_large <- function(df, large) {  
  large <- rlang::enquo(large)  
  ppr3 <- dplyr::mutate(df,  
                        is_big =  
                          ifelse(  
                            .data$price >= !!large,  
                            "Big", "Not Big"))  
  return(ppr3)  
}
```

- By which I mean copying code into a function and adding a return statement
- This is the **extract function** refactoring

# Opposite but Equal

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- Sometimes, we actually want the function to be inlined
- This is called **inline function**

# Convert price to log

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- Logs are great
- In general, if you have problems with a response or predictor variable, you should log it and see if that helps

```
ppr4 <- mutate(ppr3, log_price =  
                  log(price, base = 10))
```

- Does this need to be a function?

# Log function tests

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
test_that('log(price) < price',  
{  
  data(ppr)  
  ppr2 <- normalise_names(ppr) %>%  
    dplyr::mutate(price=fix_price(price))  
  
  new_df <- log_column(ppr2, price)  
  # can't find a vector based expectation in testthat  
  expect_lt(new_df$log_price[1],  
            ppr2$price[1]))  
}
```

- Obviously this fails again
- One of the great things about TDD is that you end up needing to use your API somewhere before you write it
- This provides both documentation and a usability check

# Log Column Function

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- Again, we need to do the pointless tidyverse dance because Hadley hates quotes

```
log_column <- function(df, col) {  
  col <- rlang::enquo(col)  
  res <- dplyr::mutate(df,  
                        log_price =  
                          log(!col, base = 10))  
  return(res)  
}
```

- This is beginning to annoy me a little less, from repetition

# What not to functionalise

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- Exploratory analysis
- There's loads of that in my scripts, for example

```
post_codes_table <-  
  with(ppr4, table(postal_code, useNA = "always"  
    as.data.frame() %>%  
    arrange(desc(Freq)))  
head(post_codes_table, n = 12)
```

- This is to figure out what's going on
- Not sure it makes a lot of sense here
- If I'm doing the same stuff on a bunch of columns, a function is useful, but otherwise it doesn't make much sense

# Next Steps

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
ppr5 <- mutate(ppr4,  
  is_full_market_price = ifelse(  
    not_full_market_price == "No",  
    "Yes",  
    "No"  
  )  
) %>%  
  select(-not_full_market_price) # remove old va
```

- So this is only done once, and is pretty specific
- But it handles some ugly, ugly names
- This is a **rename field** refactoring (which can be incredibly impactful)

# Rename Field Test

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- First, we write a test

```
data(ppr)
test_that('we have is_full_market_price column', {
  ppr3 <- normalise_names(ppr) %>%
    dplyr::mutate(price=fix_price(price)) %>%
    log_column(price)
  ppr4 <- invert_field(ppr3,
                       not_full_market_price)
  expect_equal(names(ppr4)[length(ppr4)],
               "is_full_market_price") }
)
```

- Which obviously fails
- I did reconsider my name after the failure (the original name was `fix_field_names`, but that was needlessly vague)



# Invert Field Names

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
invert_field <- function(df, field) {  
  field <- rlang::enquo(field)  
  ppr5 <- dplyr::mutate(df,  
    is_full_market_price = ifelse(  
      !!field == "No",  
      "Yes",  
      "No"  
    )  
  ) %>%  
  dplyr::select(-not_full_market_price)  
  return(ppr5)  
}
```

- Remember that once we have tests on everything, we can refactor much more fearlessly

# Description of property field

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- When I wrote this code (for a data science bootcamp) I discovered that this field is buggy and annoying
- There's also a bunch of entries as Gaeilge, which is annoying for my purposes
- I have about 30 lines of logic here, most of which won't be reusable right away.
- The easiest thing to do with it is just dump it all into a function so we can test it

# Property Description Function

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
fix_property_description <- function(df) {  
  fix_new <- mutate(df,  
    description_of_property =  
      ifelse(  
        grepl("Nua", x = description_o  
        "New Dwelling house /Apartment  
        ifelse(  
          grepl("Ath", x = descr  
          "Second-Hand Dwelling  
          ## nested ifelse are p  
          description_of_propert  
        )  
      )  
    )  
  )  
  
  remove_irish <- dplyr::mutate(fix_new, prop_description
```

# Testing our Property Description Logic

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- Kinda painful
- Let's punt on it and just ensure that the same data is output
- This is something that's a really, really common occurrence<sup>7</sup>

---

<sup>7</sup>especially with code you didn't write yourself

# Creating base data

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

Alter our `prep_modelling` script like so:

```
readr::write_csv(ppr10,  
                  path = "ppr_data_cleaning_done.csv")
```

- This will write the file into our package

# Data Splitting

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- At this point, we have created functions for each of our data cleaning steps
- The next section is handling test and train sets

# Data Splitting Test

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
data(ppr)
ppr_for_split <- normalise_names(ppr) %>%
  dplyr::mutate(price=fix_price(price)) %>%
  mark_values_as_large(1e6) %>%
  log_column(price) %>%
  invert_field(not_full_market_price) %>%
  fix_property_description()
test_that('split_data returns a list',
          expect_is(split_data(ppr_for_split), 'list'))
```

- Start with a simple test, make them more specific as times goes on

# Data Splitting Function

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
split_data <- function(df) {  
  return(list())  
}
```

- Passes the test :)



# Splitting Data

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- Dump all the code into a function
- Passes the tests

```
split_data <- function(df) {  
  set.seed(34)  
  ppr_train_indices <- with(  
    df,  
    caret::createDataPartition(log_price,  
                                times = 1,  
                                p = 0.7,  
                                list = FALSE  
  )  
  ) %>% as.vector() #because tibble sucks  
  ppr_train <- df[ppr_train_indices, ]  
  
  ppr_not_train <- df[-ppr_train_indices, ]  
  ppr_test_indices <- with(  
    df,
```

# Add More Tests

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
test_that('split data has test and train',  
          expect_equal(names(  
            split_data(ppr_for_split)),  
            c("train", "test"))))  
  
test_that('split data returns a train tibble',  
          expect_is(  
            split_data(ppr_for_split)$train[1],  
            'tbl_df'))
```

# Yet More Tests

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
test_that('split data returns a test tibble',
  expect_is(
    split_data(ppr_for_split)$test[1],
    'tbl_df'))

test_that('split_data train has less rows than input',
  train <- split_data(ppr_for_split)$train
  expect_gt(
    nrow(ppr_for_split),
    nrow(train))
)
```

# Feature Engineering

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- We now move on to the next script
- First thing we do is add a `write_csv` call to the end of our original function

```
readr::write_csv(test_vtreat_full,  
                  path="feature_eng_results_old.csv")
```

# Feature Engineering Integration Test

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
context("integration test feature engineering")
old <- readr::read_csv("../../inst/feature_eng_results")
new <- readr::read_csv("../../inst/feature_eng_results")
test_that(
  "data.frame outputs are equal",
  expect_equal(old, new)
)
```

# Analysing what the script does

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- Takes the data from the previous script
- fits a simple glmnet model
- loads in geocoded register
- Loads in more data (ppr plus pobal deprivation data)
- fits a PCA
- fit a linear model on electoral district
- loads vtreat
- applies vtreat functionality to df
- outputs model matrix <sup>8</sup> from vtreat

---

<sup>8</sup>it's actually a df, but it is very very similar to a `model.matrix`

# Cleaning it up

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- We'll focus on the pathway to the vtreat model matrix
- Ignore the PCA and the simple models for now

# Create More Datasets

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- We need to load in a sample of our geocoded and pobal data
- Again, this makes it way easier and more reproducible to run the tests

```
ppr_gc <- read_csv(  
  "~/Dropbox/PPR/ppr_geocoded_till_oct2018.csv"  
)  
set.seed(49)  
ppr_sample <- dplyr::sample_frac(ppr_gc, size = 0.05)  
ppr_gc <- ppr_sample  
usethis::use_data(ppr_gc, overwrite = TRUE)
```



# Pobal Data

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
ppr_pobal <- readRDS("~/Dropbox/Code/DDS/ppr_sf_pobal2")
set.seed(49)
ppr_pobal_sample <- dplyr::sample_frac(ppr_pobal, size = 0.1)
ppr_pobal <- ppr_pobal_sample
usethis::use_data(ppr_pobal, overwrite = TRUE)
```

# Run these scripts

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- Weirdly, despite these being in the package directory, nothing is done unless you execute the scripts into a running process
- So make sure you do that
- You also need to document your datasets, but we'll avoid that for now
- Note that you'll need to update the integration test after adding data to the package
- Also note that sampling breaks the integration tests, unless you use `set.seed`

# Write Some Tests

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

# Vtreat Pipeline

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
generate_model_matrix <- function(df, varlist,
                                   outcomename,
                                   calibration_size) {
  df_calibration <- dplyr::sample_frac(df, size=calibration_size)
  df_train <- dplyr::anti_join(df, df_calibration)
  tf <- designTreatmentsN(df,
                          varlist = varlist,
                          outcomename = outcomename
                          )

  result <- prepare(tf,
                    dframe = df_train
                    )

  return(result)
}
```

# Subtleties and Problems

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- Sampling will break lots of tests given the approach we're using
- You can either `set.seed` or reduce the input data outside

# Different Approach to Refactoring

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- Instead of abstracting out units of work, abstract out common tasks

```
sapply(ppr_pobal, n_distinct) %>%  
  as.data.frame() %>%  
  rownames_to_column() %>%  
  arrange(desc('.', '.'))
```

- This code is unique
- But the pattern is common
- This is a lot more useful than the original
- It's also a lot more amenable to testing
- As always, we start with a characterisation test

# More Extractions

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
get_max_and_min <- select_if(as.data.frame(ppr),  
                             is.numeric) %>%  
  sapply(., function(x) {  
    data.frame(  
      min = min(x, na.rm = TRUE)  
      max = max(x, na.rm = TRUE)  
    )  
  }) %>%  
as.data.frame()
```

# Becomes

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
get_max_and_min <- function(df) {  
  maxmin <- dplyr::select_if(as.data.frame(df),  
                             is.numeric) %>%  
    sapply(., function(x) {  
      data.frame(  
        min = min(x, na.rm = TRUE),  
        max = max(x, na.rm = TRUE)  
      )  
    }) %>%  
    tibble::as_tibble() %>%  
    tibble::rownames_to_column()  
  
  max_min_df <- tidyr::unnest(maxmin,  
                             cols=colnames(maxmin)  
                             as.data.frame())  
  return(max_min_df)
```



# Value Counts

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- This is *literally* the only thing pandas does better than R
- I write code like the below far too often for my taste

```
post_codes_table <-  
  with(ppr4, table(postal_code,  
                    useNA = "always")) %>%  
    as.data.frame() %>%  
    arrange(desc(Freq))
```

# Which can be functionalised like so

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
value_counts <- function(df, col) {  
  result <-  
    table(eval(substitute(col), envir=df),  
          useNA = "always") %>%  
    as.data.frame() %>%  
    arrange(desc(Freq))  
  
}
```

# Proportion of Missings

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- Again, I have written code like this many, many times

```
count_proportion_missing <- function(df) {  
  result <- sapply(df, function(x) {  
    sum(is.na(x)) / length(x)  
  }) %>%  
    as.data.frame() %>%  
    tibble::rownames_to_column() %>%  
    dplyr::arrange(desc('.')) %>%  
    dplyr::mutate_if(is.numeric, round, 4)  
  return(result)  
}
```

# I don't actually need these functions to run the script

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- But I can add the data from the originals as test data to get started
- And thus ensure that my changes don't break anything
- To be honest, the last few functions are probably the most useful thing in my package right now
- Certainly, they are the most re-usable
- They do need some tests

# Lazy Man's approach to testing

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- This is my philosophy
- Minimum testing for maximum output
- We add a few `write_csv` calls for each of our outputs
- And use that data for testing

# Write a function to make this easier

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- We have lots of annoying relative paths in our test data loading
- let's abstract this away from our concerns

```
load_test_data <- function(name) {  
  bpath <- "../../inst/"  
  result <- readr::read_csv(paste0(bpath, name))  
  return(result)  
}
```

# Tests for our Data Analysis Functions

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
require(ppr)
require(testthat)
count_distincts_old <- load_test_data("count_distincts")
test_that('count_distinct old and new are equal',
          expect_equal(count_distinct_values(ppr_pobal,
                                             count_distincts_old))
```

- And it fails <sup>9</sup>

---

<sup>9</sup>always make sure your tests fail when they're supposed to

# More Tests

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
maxmin_old <- load_test_data("get_max_and_min_test_data")
test_that('get_max_and_min old and new are equal',
  expect_equal(get_max_and_min(ppr_pobal),
    maxmin_old))
```



# Yet More Tests

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

```
missings_old <- load_test_data("count_proportion_missi  
test_that('count_prop_missings old and new are equal',  
          expect_equal(count_proportion_missing(ppr_po  
                        missings_old))
```

# Further Work

Seeing the  
Light with  
Refactoring  
and Testing

Richie  
Morrisroe

- I haven't covered that much refactoring
- I haven't covered that much TDD
- I haven't covered that much working with legacy code
- But hopefully I have given an introduction to them